Chapter 13

Simple Linear
Regression

Learning Objectives

In this chapter, you learn:

- How to use regression analysis to predict the value of a dependent variable based on a value of an independent variable
- The meaning of the regression coefficients b_0 and b_1
- How to evaluate the assumptions of regression analysis and know what to do if the assumptions are violated
- To make inferences about the slope and correlation coefficient
- To estimate mean values and predict individual values

Correlation vs. Regression

DCOVA

- A scatter plot can be used to show the relationship between two variables
- Correlation analysis is used to measure the strength of the association (linear relationship) between two variables
  - Correlation is only concerned with strength of the relationship
  - No causal effect is implied with correlation

Types of Relationships

DCOVA

Linear relationships

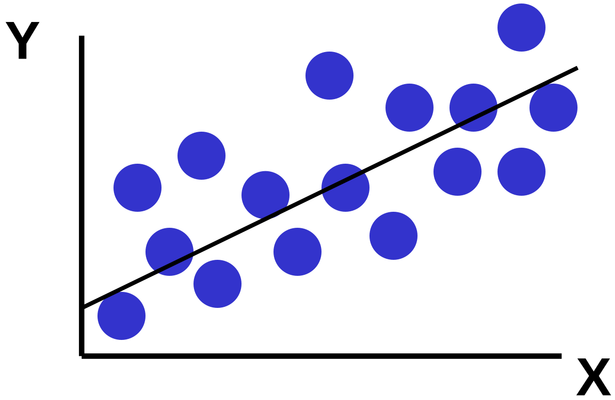

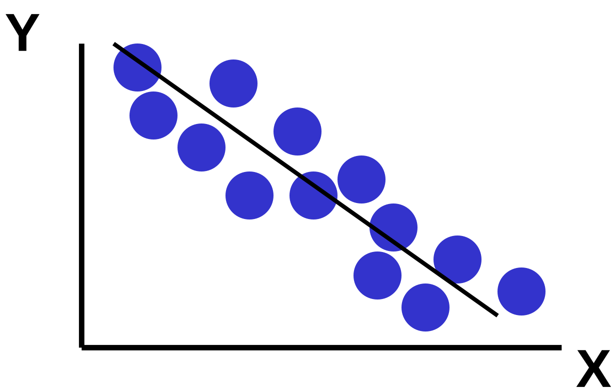

Curvilinear relationships

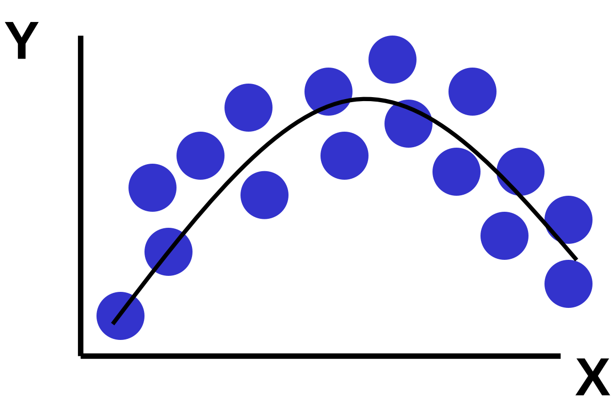

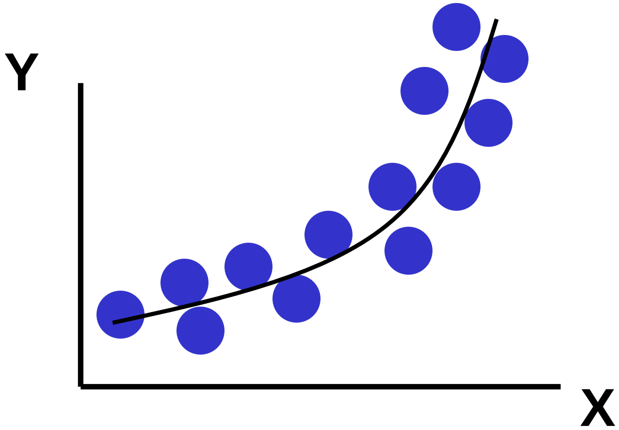

Types of Relationships

DCOVA

(continued)

Strong relationships

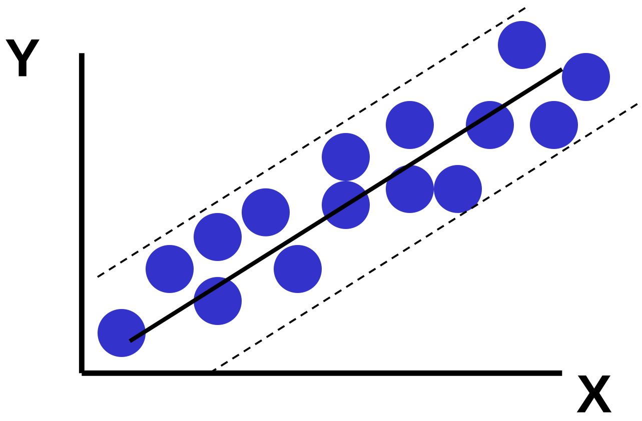

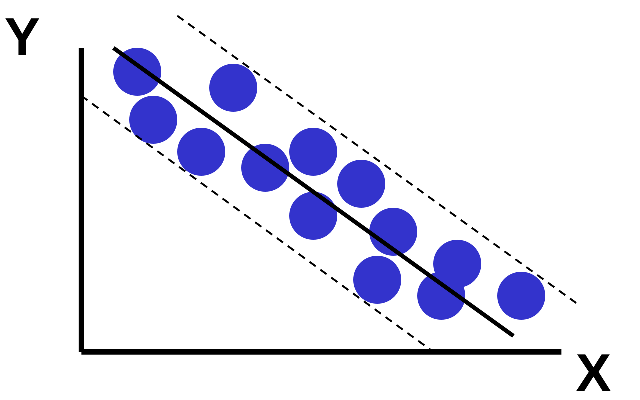

Weak relationships

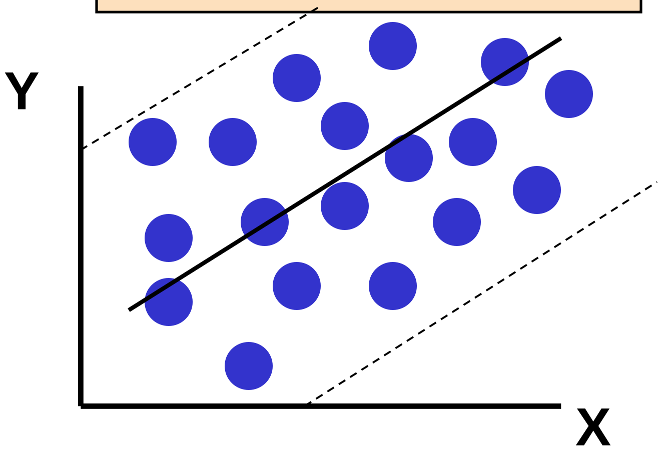

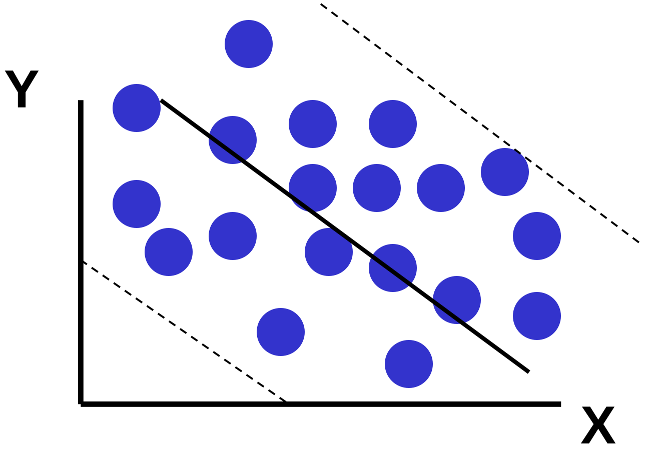

Types of Relationships

DCOVA

(continued)

No relationship

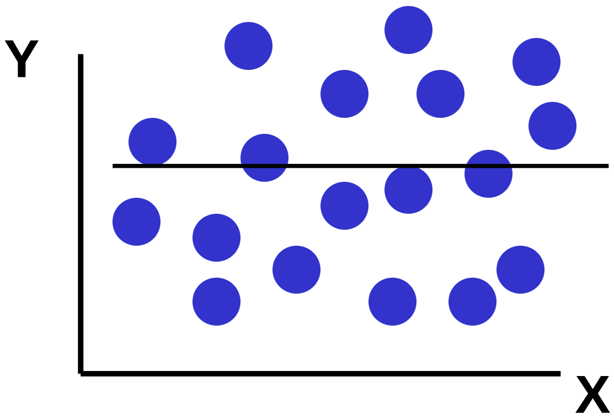

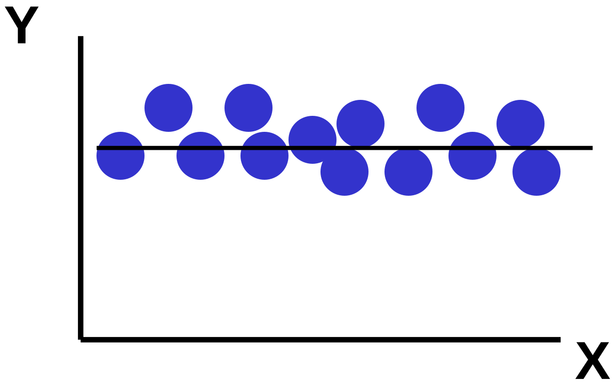

Introduction to Regression Analysis

DCOVA

- Regression analysis is used to:
  - Predict the value of a dependent variable based on the value of at least one independent variable
  - Explain the impact of changes in an independent variable on the dependent variable

Dependent variable: the variable we wish to predict or explain

Independent variable: the variable used to predict or explain the dependent variable

Simple Linear Regression Model

DCOVA

- Only **one** independent variable, X
- Relationship between X and Y is described by a linear function
- Changes in Y are assumed to be related to changes in X

Simple Linear Regression Model

DCOVA

Population Y intercept

Population Slope Coefficient

Independent Variable

Random Error term

Dependent Variable

$$Y_i = \underbrace{\beta_0 + \beta_1 X_i}_{\text{Linear component}} + \underbrace{\varepsilon_i}_{\text{Random Error component}}$$

Simple Linear Regression Model

DCOVA

(continued)

$$Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$$

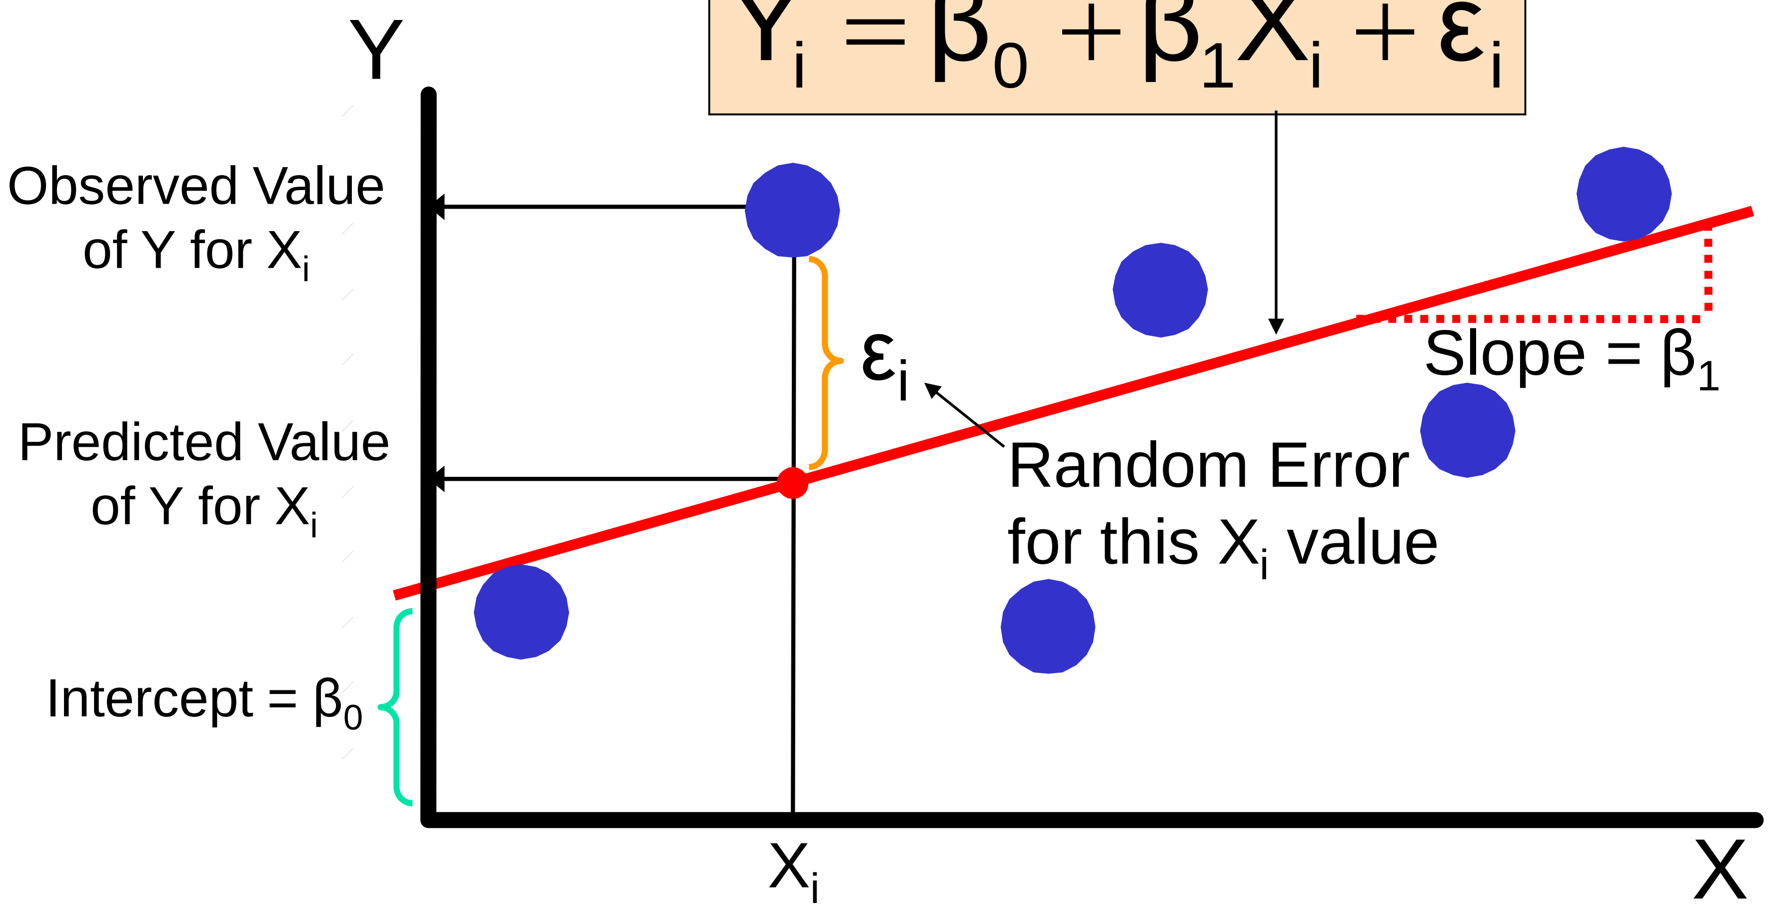

Simple Linear Regression Equation (Prediction Line)

DCOVA

The simple linear regression equation provides an estimate of the population regression line

Estimated (or predicted) Y value for observation i

Estimate of the regression intercept

Estimate of the regression slope

Value of X for observation i

$$\hat{Y}_i = b_0 + b_1 X_i$$

The Least Squares Method

DCOVA

b_0 and b_1 are obtained by finding the values of that minimize the sum of the squared differences between Y and $\hat{Y}$:

$$\min \sum (Y_i - \hat{Y}_i)^2 = \min \sum (Y_i - (b_0 + b_1 X_i))^2$$

Finding the Least Squares Equation

DCOVA

- The coefficients b_0 and b_1 , and other regression results in this chapter, will be found using Excel or Minitab

Formulas are shown in the text for those who are interested

Interpretation of the Slope and the Intercept

DCOVA

- b_0 is the estimated average value of Y when the value of X is zero
- b_1 is the estimated change in the average value of Y as a result of a one-unit increase in X

Simple Linear Regression Example

DCOVA

- A real estate agent wishes to examine the relationship between the selling price of a home and its size (measured in square feet)

- A random sample of 10 houses is selected
  - Dependent variable (Y) = house price in $1000s
  - Independent variable (X) = square feet

Simple Linear Regression Example: Data

DCOVA

House Price in $1000s (Y)	Square Feet (X)
245	1400
312	1600
279	1700
308	1875
199	1100
219	1550
405	2350
324	2450
319	1425
255	1700

Simple Linear Regression Example: Scatter Plot

DCOVA

House price model: Scatter Plot

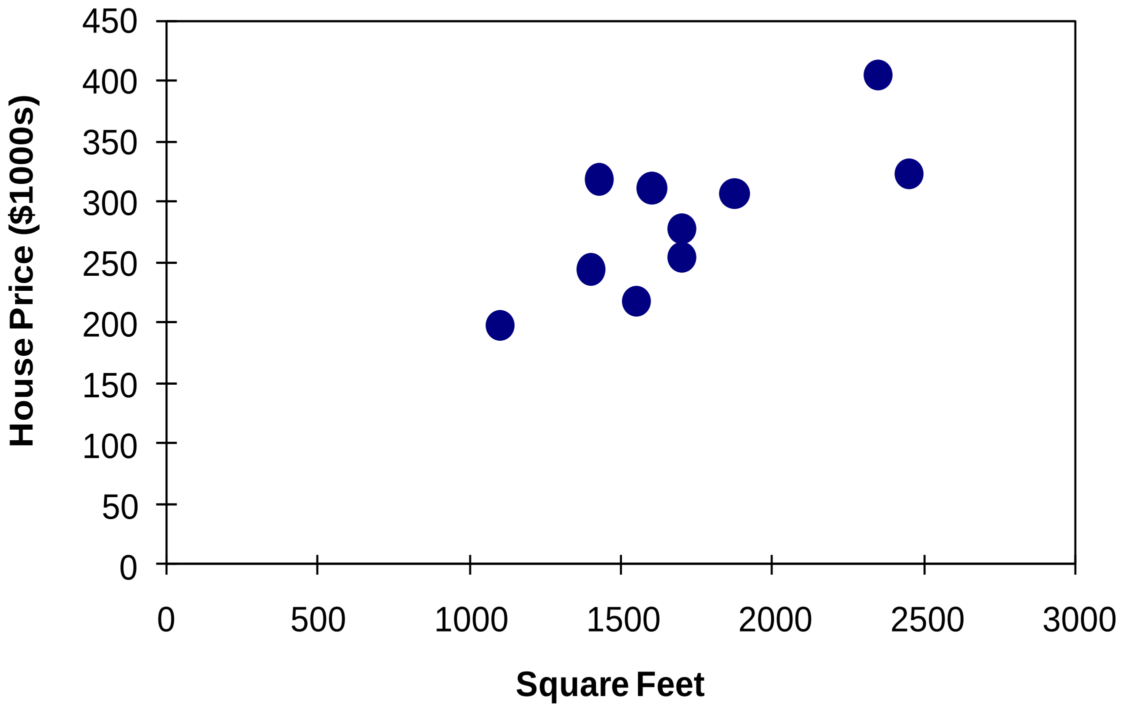

Simple Linear Regression Example: Minitab Output

DCOVA

The regression equation is

Price = 98.2 + 0.110 Square Feet

Predictor	Coef	SE Coef	T	P
Constant	98.25	58.03	1.69	0.129
Square Feet	0.10977	0.03297	3.33	0.010

S = 41.3303   R-Sq = 58.1%   R-Sq(adj) = 52.8%

Analysis of Variance

Source	DF	SS	MS	F	P
Regression	1	18935	18935	11.08	0.010
Residual Error	8	13666	1708		
Total	9	32600			

The regression equation is:

$$\widehat{\text{house price}} = 98.24833 + 0.10977\,(\text{square feet})$$

Simple Linear Regression Example: Graphical Representation

DCOVA

House price model: Scatter Plot and Prediction Line

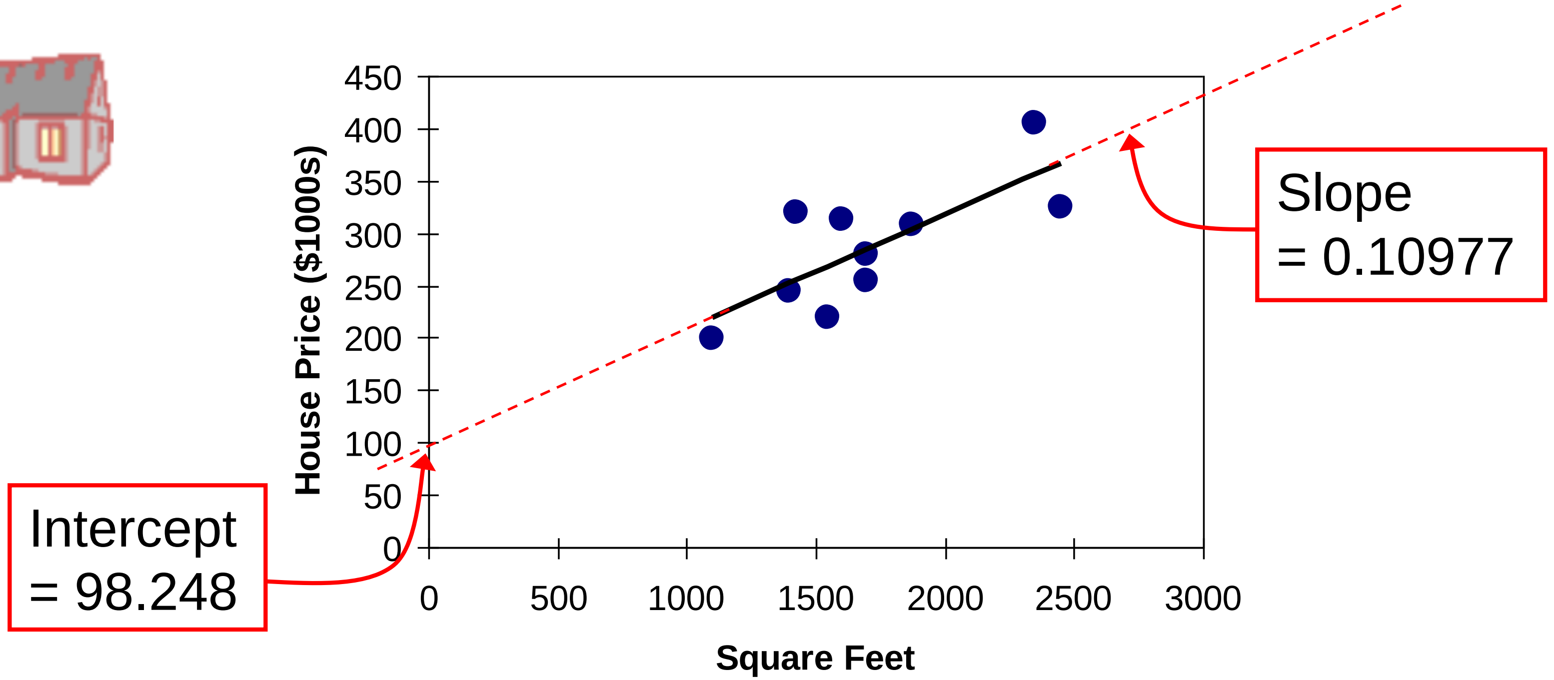

$$\widehat{\text{house price}} = 98.24833 + 0.10977\ (\text{square feet})$$

Simple Linear Regression Example: Interpretation of b_o

DCOVA

$$\widehat{\text{house price}} = \boxed{98.24833} + 0.10977\ (\text{square feet})$$

- b_0 is the estimated average value of Y when the value of X is zero (if X = 0 is in the range of observed X values)
- Because a house cannot have a square footage of 0, b_0 has no practical application

Simple Linear Regression Example: Interpreting b_1

DCOVA

$$\widehat{\text{house price}} = 98.24833 + \boxed{0.10977}\,(\text{square feet})$$

- b_1 estimates the change in the average value of Y as a result of a one-unit increase in X
  - Here, $\boxed{b_1 = 0.10977}$ tells us that the mean value of a house increases by .10977($1000) = $109.77, on average, for each additional one square foot of size

Simple Linear Regression Example: Making Predictions

DCOVA

Predict the price for a house with 2000 square feet:

$$\widehat{\text{house price}} = 98.25 + 0.1098\ (\text{sq.ft.})$$

$$= 98.25 + 0.1098(2000)$$

$$= 317.85$$

The predicted price for a house with 2000 square feet is 317.85($1,000s) = $317,850

Simple Linear Regression Example: Making Predictions

DCOVA

- When using a regression model for prediction, only predict within the relevant range of data

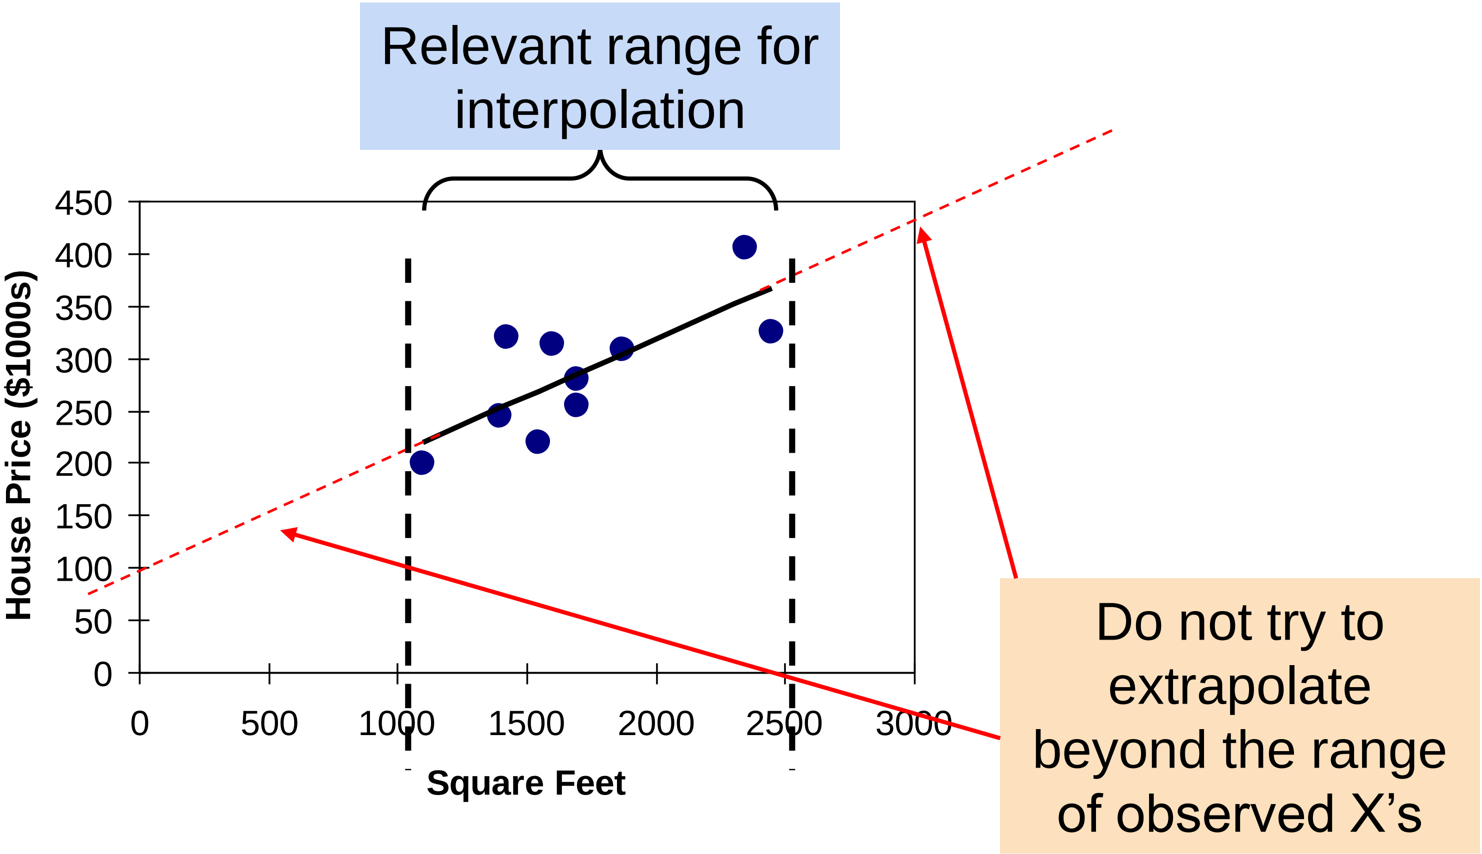

Measures of Variation

DCOVA

- Total variation is made up of two parts:

$$SST = SSR + SSE$$

Total Sum of Squares	Regression Sum of Squares	Error Sum of Squares

$$SST = \sum (Y_i - \bar{Y})^2 \qquad SSR = \sum (\hat{Y}_i - \bar{Y})^2 \qquad SSE = \sum (Y_i - \hat{Y}_i)^2$$

where:

$\bar{Y}$ = Mean value of the dependent variable

Y_i = Observed value of the dependent variable

$\hat{Y}_i$ = Predicted value of Y for the given X_i value

Measures of Variation

(continued)

DCOVA

- SST = total sum of squares (Total Variation)
  - Measures the variation of the Y_i values around their mean $\bar{Y}$
- SSR = regression sum of squares (Explained Variation)
  - Variation attributable to the relationship between X and Y
- SSE = error sum of squares (Unexplained Variation)
  - Variation in Y attributable to factors other than X

Measures of Variation

(continued)

DCOVA

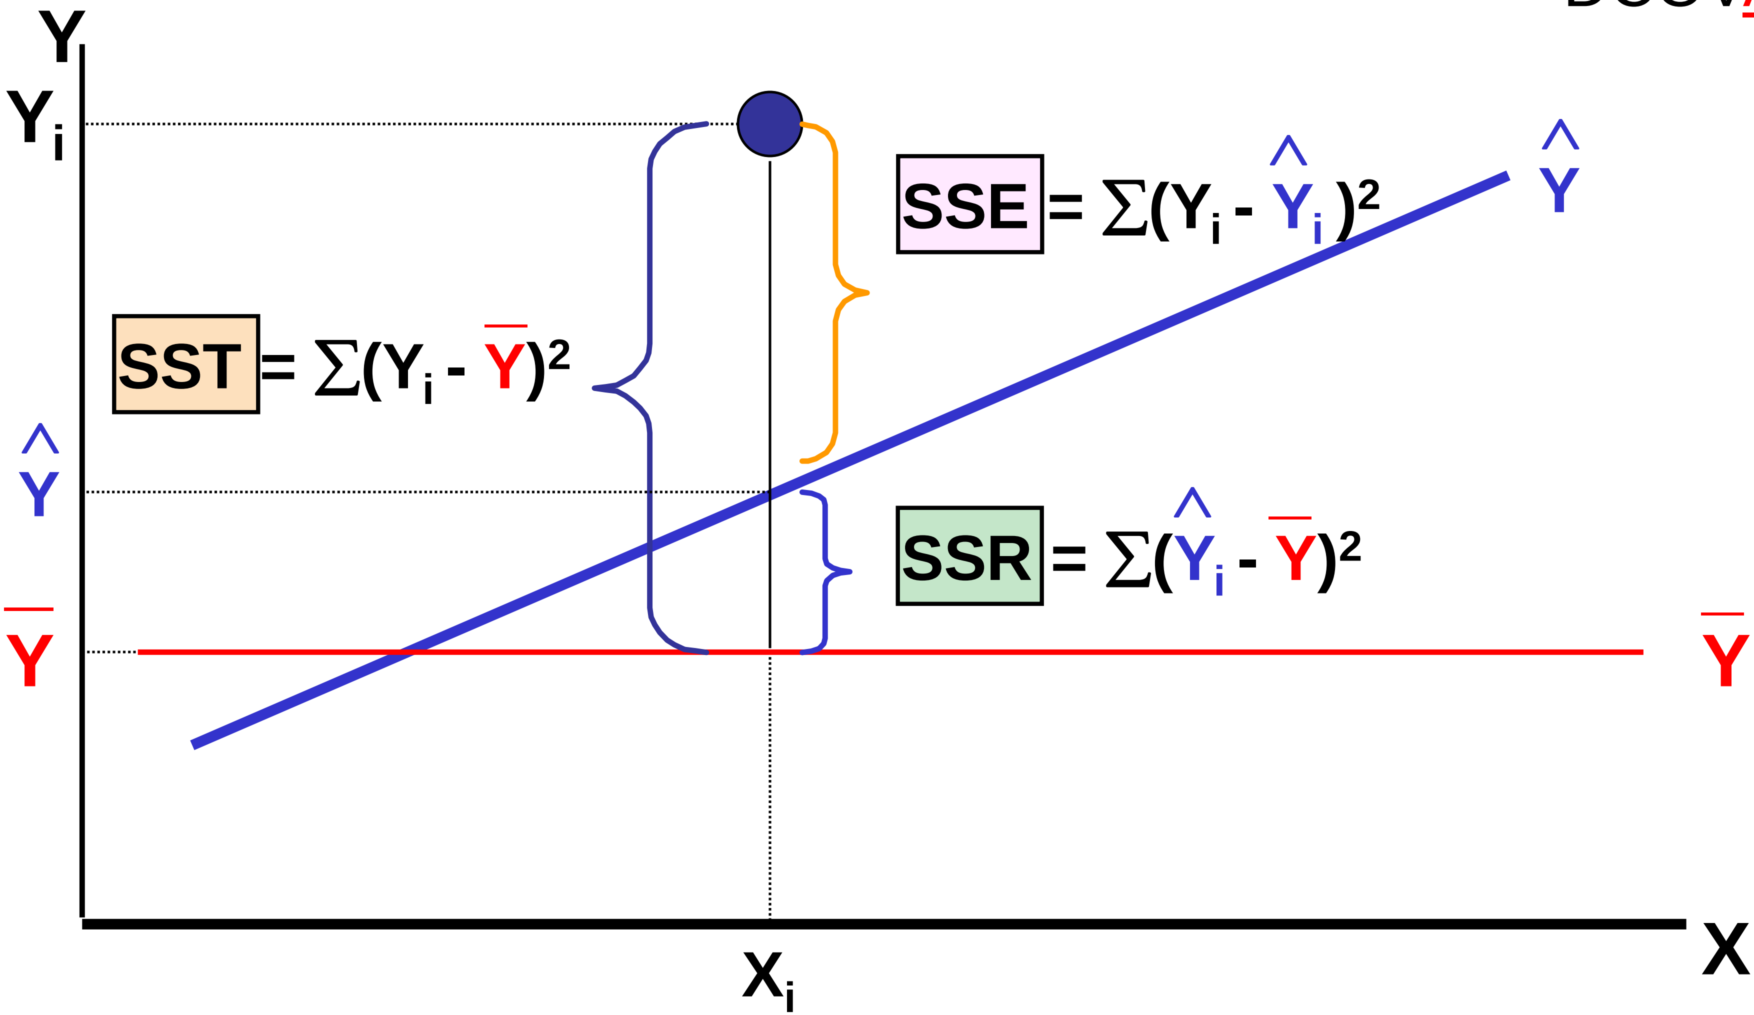

Coefficient of Determination, r^2

DCOVA

- The coefficient of determination is the portion of the total variation in the dependent variable that is explained by variation in the independent variable
- The coefficient of determination is also called r-squared and is denoted as r^2

$$r^2 = \frac{SSR}{SST} = \frac{\text{regression sum of squares}}{\text{total sum of squares}}$$

note: $0 \le r^2 \le 1$

Examples of Approximate r^2 Values

DCOVA

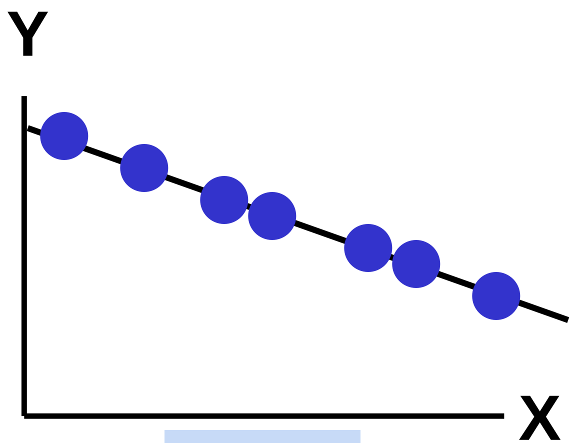

$r^2 = 1$

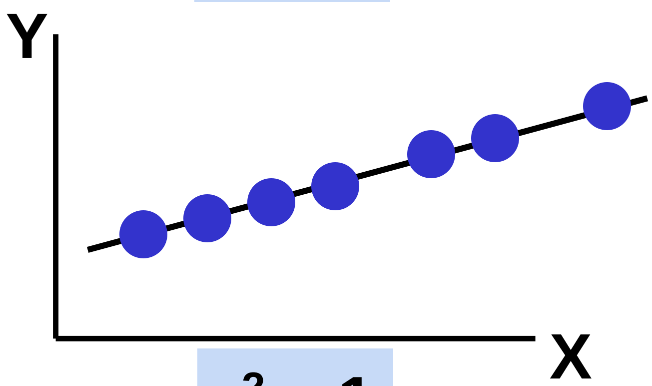

$r^2 = 1$

$r^2 = 1$

Perfect linear relationship between X and Y:

100% of the variation in Y is explained by variation in X

Examples of Approximate r^2 Values

DCOVA

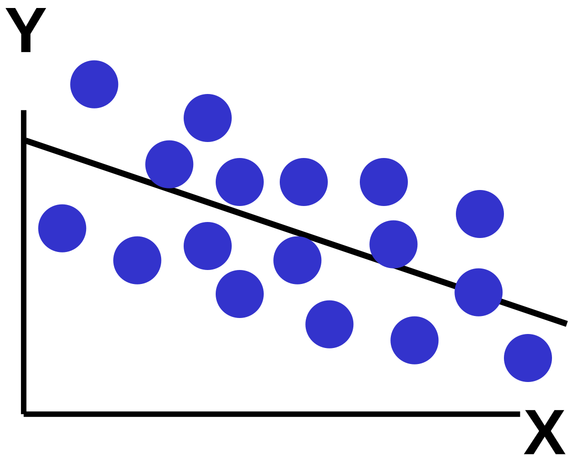

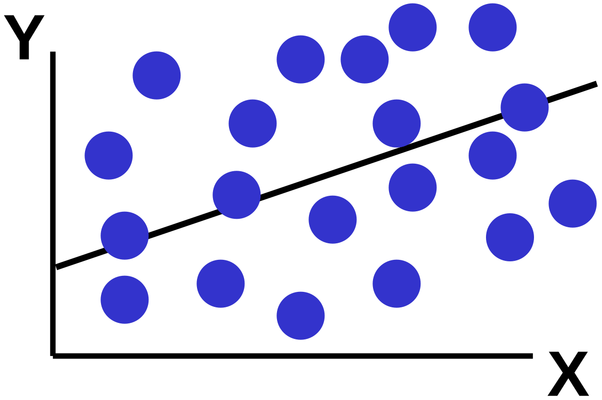

$0 < r^2 < 1$

Weaker linear relationships between X and Y:

Some but not all of the variation in Y is explained by variation in X

Examples of Approximate r^2 Values

DCOVA

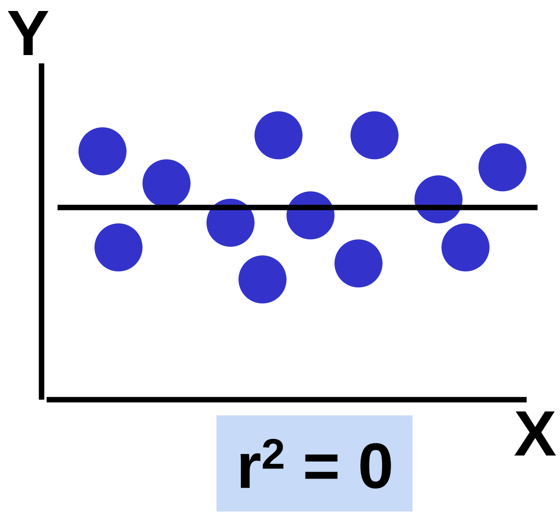

$r^2 = 0$

No linear relationship between X and Y:

The value of Y does not depend on X. (None of the variation in Y is explained by variation in X)

Simple Linear Regression Example: Coefficient of Determination, r^2 in Minitab

DCOVA

The regression equation is

Price = 98.2 + 0.110 Square Feet

Predictor	Coef	SE Coef	T	P
Constant	98.25	58.03	1.69	0.129
Square Feet	0.10977	0.03297	3.33	0.010

S = 41.3303 R-Sq = 58.1% R-Sq(adj) = 52.8%

Analysis of Variance

Source	DF	SS	MS	F	P
Regression	1	18935	18935	11.08	0.010
Residual Error	8	13666	1708		
Total	9	32600			

$$r^2 = \frac{SSR}{SST} = \frac{18934.9348}{32600.5000} = 0.58082$$

58.08% of the variation in house prices is explained by variation in square feet

Standard Error of Estimate

DCOVA

- The standard deviation of the variation of observations around the regression line is estimated by

$$S_{YX} = \sqrt{\frac{SSE}{n-2}} = \sqrt{\frac{\sum_{i=1}^{n}(Y_i - \hat{Y}_i)^2}{n-2}}$$

Where

SSE = error sum of squares

n = sample size

Simple Linear Regression Example: Standard Error of Estimate in Minitab

DCOVA

The regression equation is

Price = 98.2 + 0.110 Square Feet

Predictor	Coef	SE Coef	T	P
Constant	98.25	58.03	1.69	0.129
Square Feet	0.10977	0.03297	3.33	0.010

S = 41.3303   R-Sq = 58.1%   R-Sq(adj) = 52.8%

Analysis of Variance

Source	DF	SS	MS	F	P
Regression	1	18935	18935	11.08	0.010
Residual Error	8	13666	1708		
Total	9	32600			

$$S_{YX} = 41.33032$$

Comparing Standard Errors

DCOVA

S_{YX} is a measure of the variation of observed Y values from the regression line

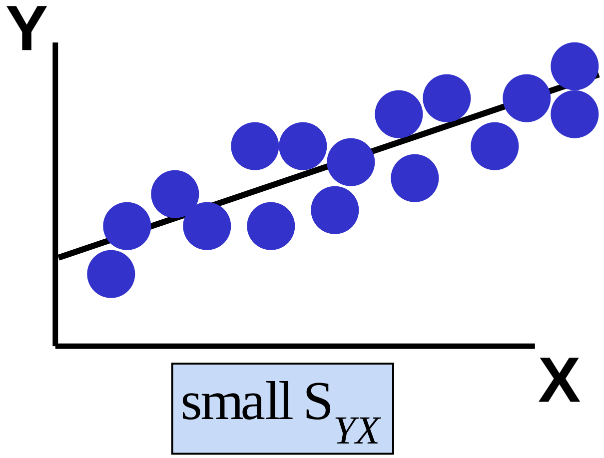

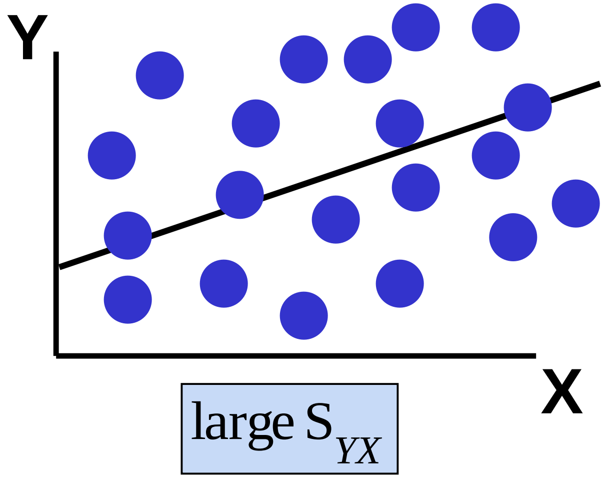

The magnitude of S_{YX} should always be judged relative to the size of the Y values in the sample data

i.e., S_{YX} = $41.33K is moderately small relative to house prices in the $200K - $400K range

Assumptions of Regression L.I.N.E

DCOVA

- Linearity
  - The relationship between X and Y is linear
- Independence of Errors
  - Error values are statistically independent
- Normality of Error
  - Error values are normally distributed for any given value of X
- Equal Variance (also called homoscedasticity)
  - The probability distribution of the errors has constant variance

Residual Analysis

DCOVA

$$e_i = Y_i - \hat{Y}_i$$

- The residual for observation i, e_i, is the difference between its observed and predicted value
- Check the assumptions of regression by examining the residuals
  - Examine for linearity assumption
  - Evaluate independence assumption
  - Evaluate normal distribution assumption
  - Examine for constant variance for all levels of X (homoscedasticity)
- Graphical Analysis of Residuals
  - Can plot residuals vs. X

Residual Analysis for Linearity

DCOVA

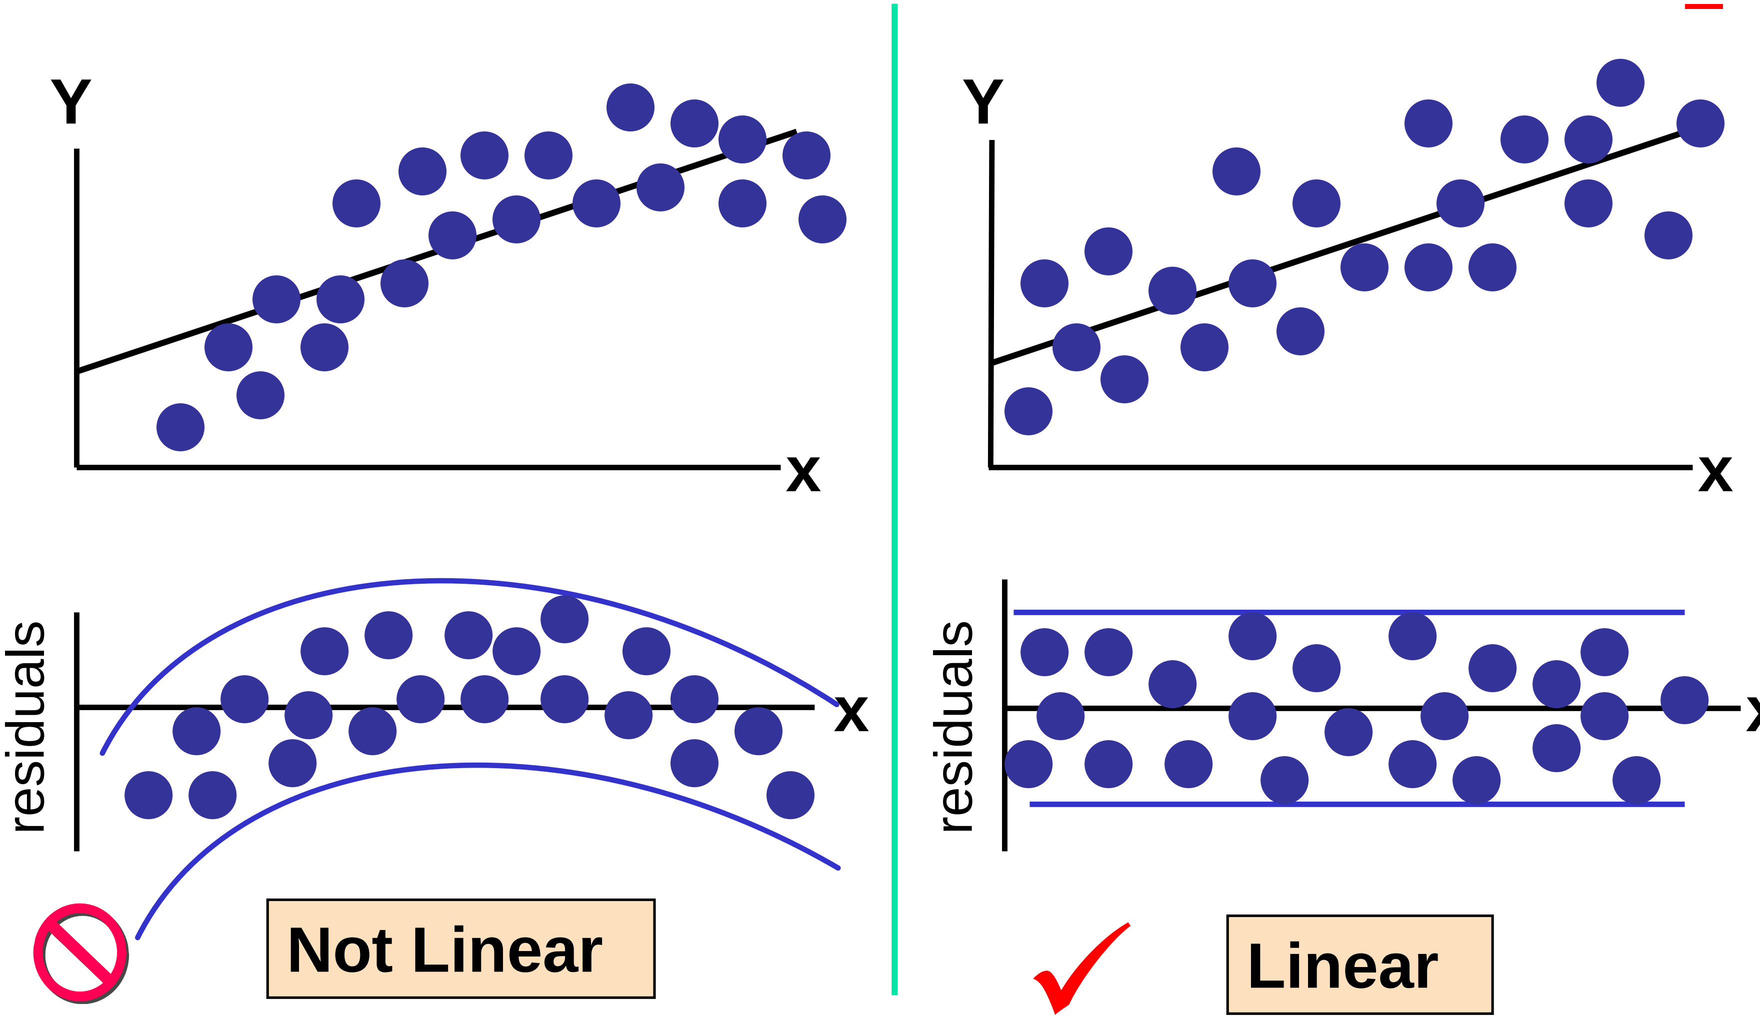

Residual Analysis for Independence

DCOVA

Not Independent

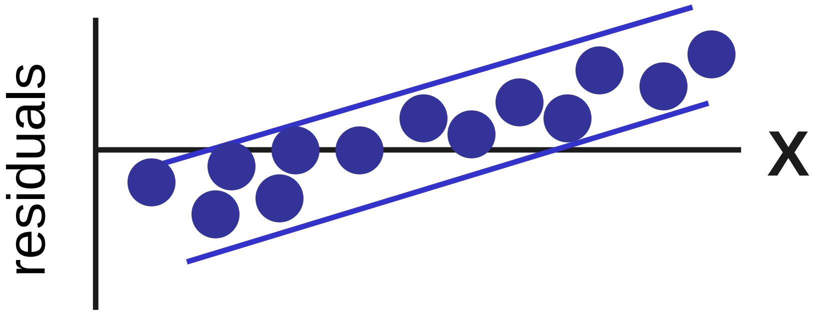

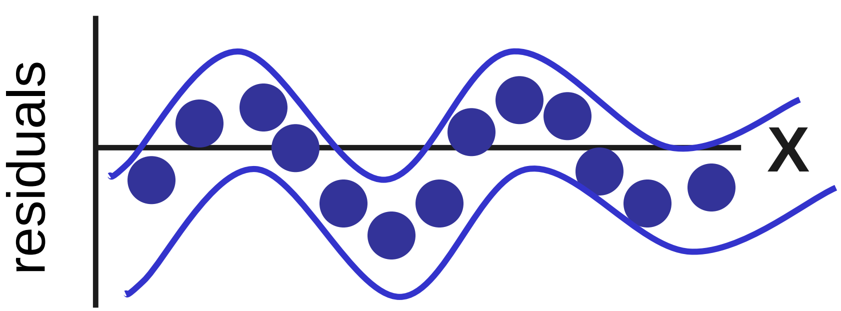

Independent

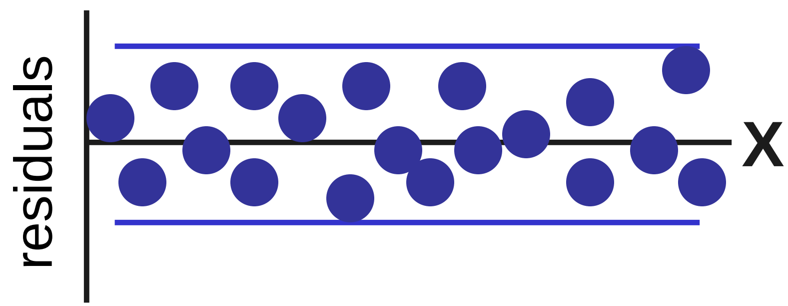

Checking for Normality

DCOVA

- Examine the Stem-and-Leaf Display of the Residuals
- Examine the Boxplot of the Residuals
- Examine the Histogram of the Residuals
- Construct a Normal Probability Plot of the Residuals

Residual Analysis for Normality

DCOVA

When using a normal probability plot, normal errors will approximately display in a straight line

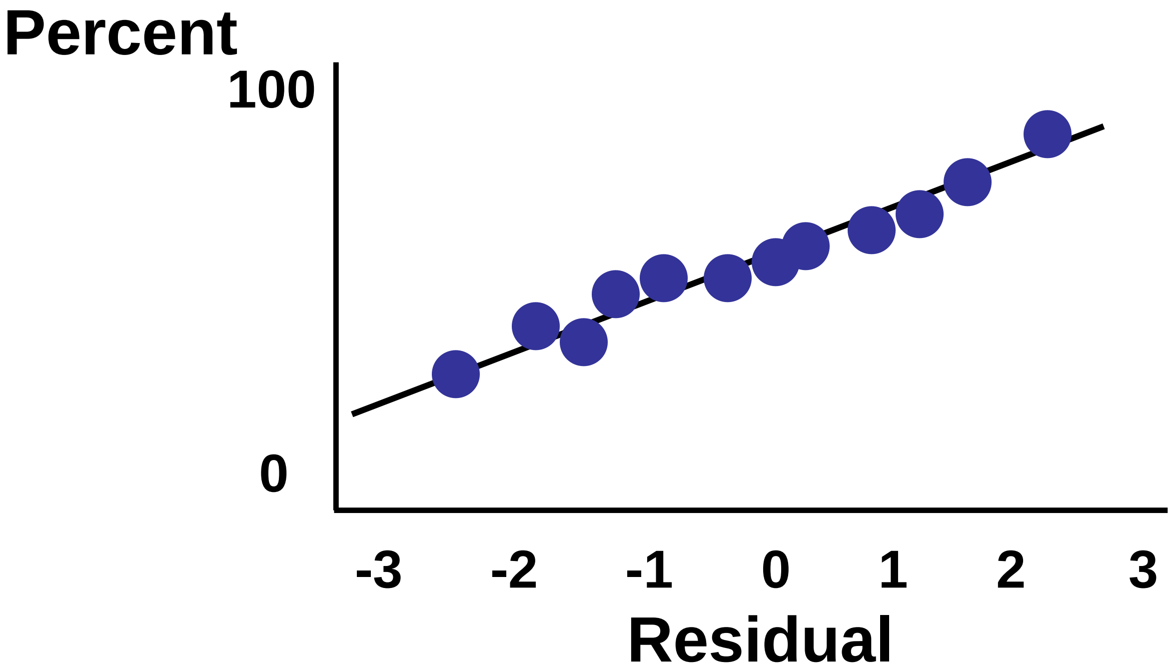

Residual Analysis for Equal Variance

DCOVA

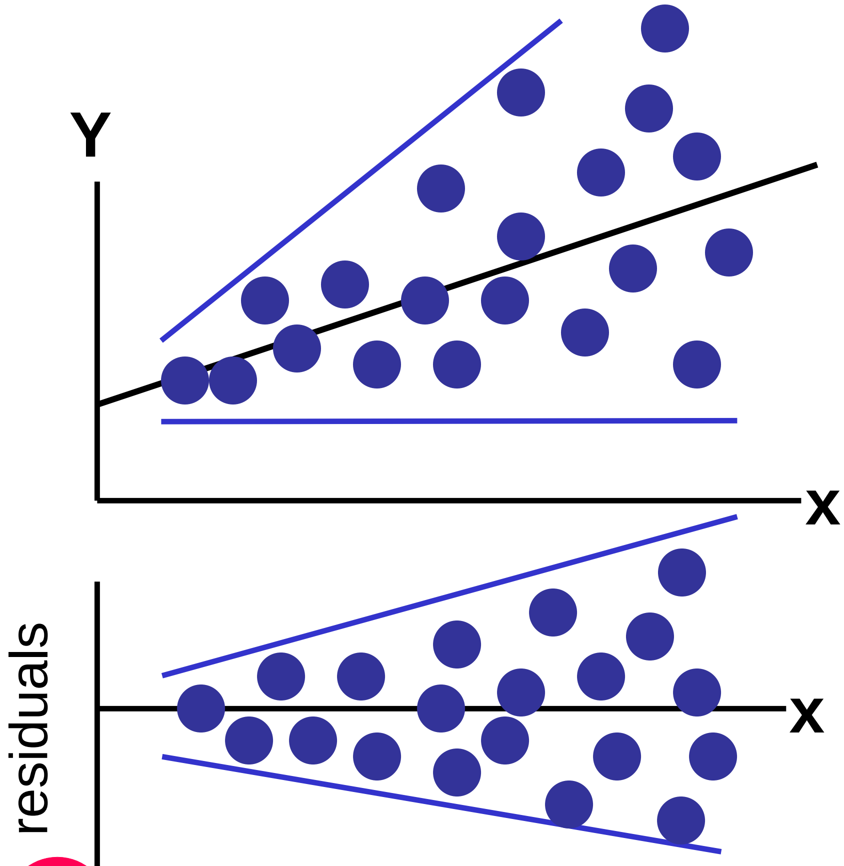

Non-constant variance

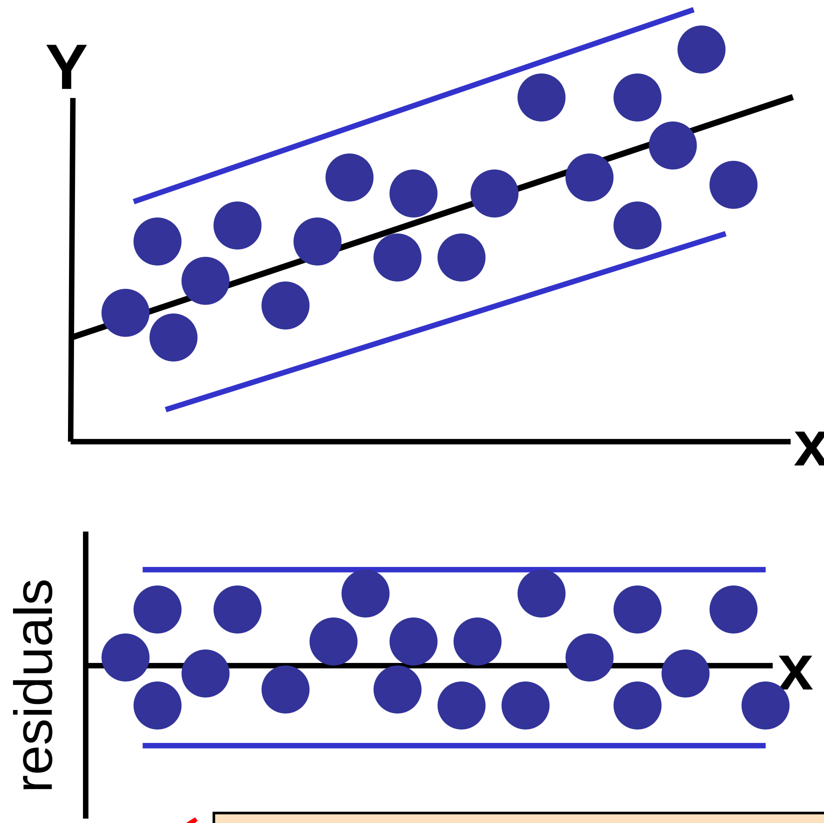

Constant variance

Simple Linear Regression Example: Minitab Residual Output

DCOVA

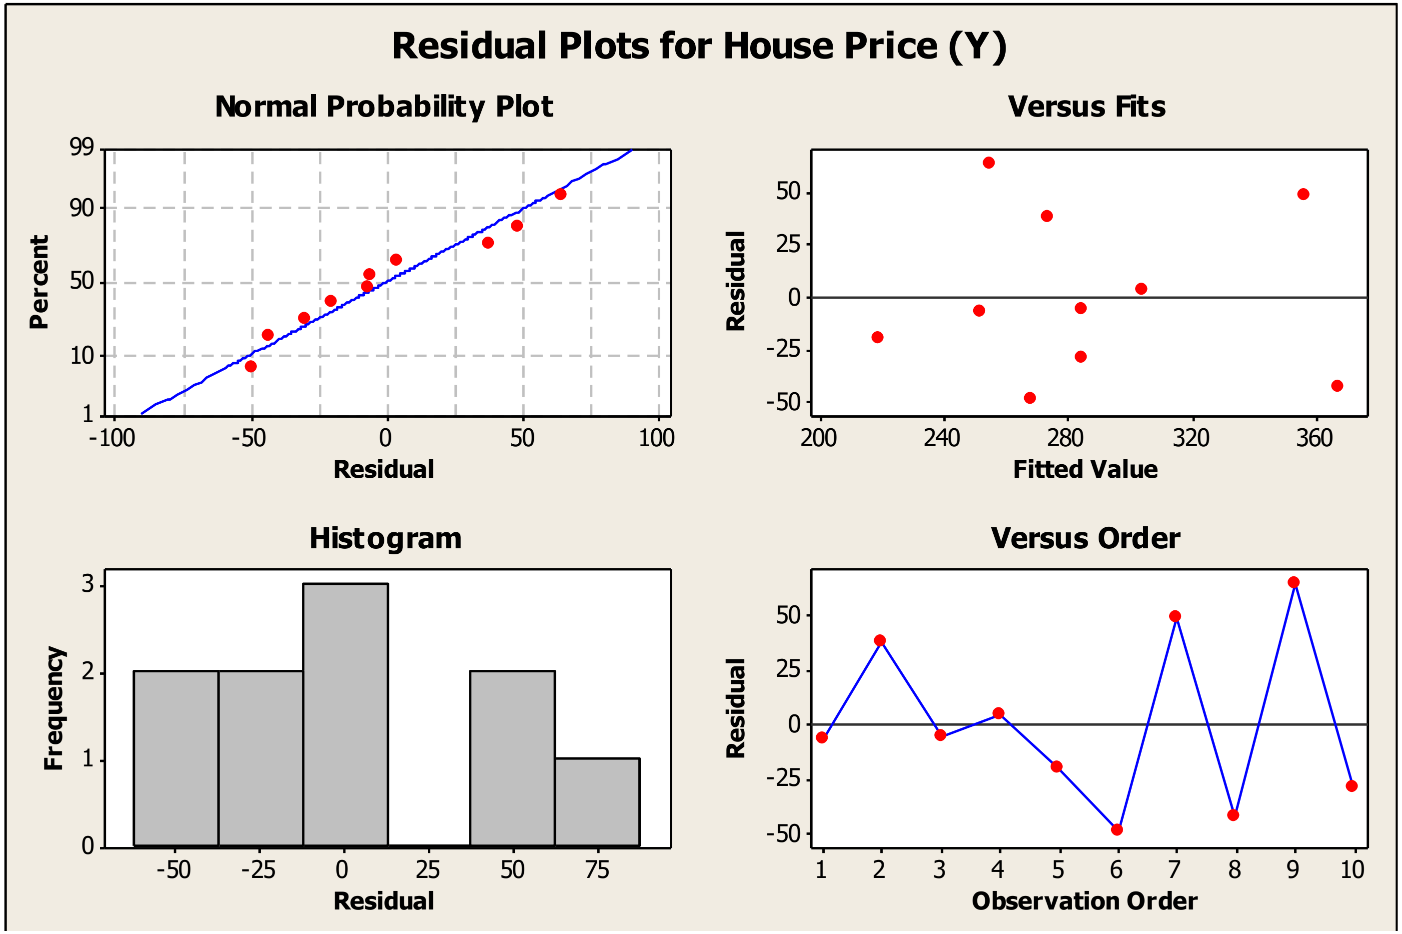

Does not appear to violate any regression assumptions

Measuring Autocorrelation: The Durbin-Watson Statistic

DCOVA

- Used when data are collected over time to detect if autocorrelation is present
- Autocorrelation exists if residuals in one time period are related to residuals in another period

Autocorrelation

DCOVA

- Autocorrelation is correlation of the errors (residuals) over time

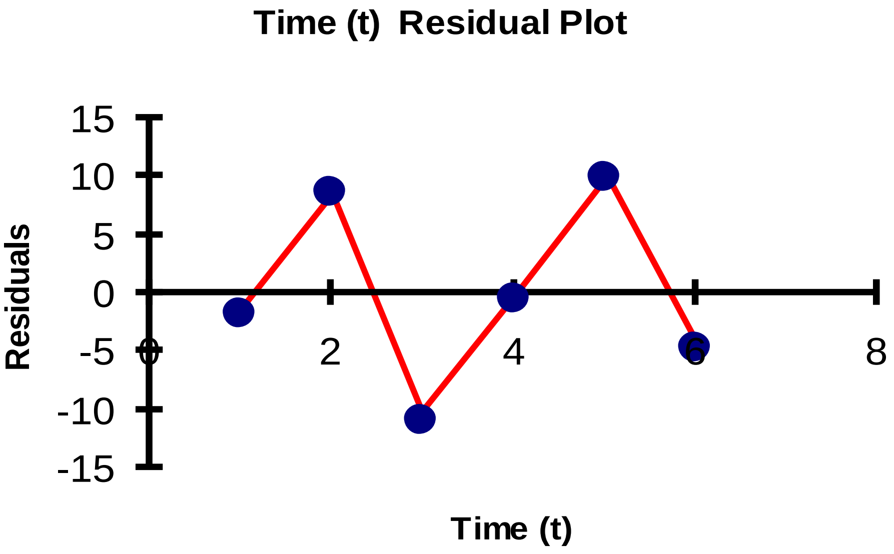

- Here, residuals show a cyclic pattern, not random. Cyclical patterns are a sign of positive autocorrelation

- Violates the regression assumption that residuals are random and independent

The Durbin-Watson Statistic

DCOVA

- The Durbin-Watson statistic is used to test for autocorrelation

H_0: residuals are not correlated
H_1: positive autocorrelation is present

$$D = \frac{\sum_{i=2}^{n}(e_i - e_{i-1})^2}{\sum_{i=1}^{n} e_i^2}$$

- The possible range is $0 \le D \le 4$
- D should be close to 2 if H_0 is true
- D less than 2 may signal positive autocorrelation, D greater than 2 may signal negative autocorrelation

Testing for Positive Autocorrelation

DCOVA

H_0: positive autocorrelation does not exist

H_1: positive autocorrelation is present

- Calculate the Durbin-Watson test statistic = D
  (The Durbin-Watson Statistic can be found using Excel or Minitab)
- Find the values d_L and d_U from the Durbin-Watson table
  (for sample size **n** and number of independent variables **k**)

Decision rule: reject H_0 if $D < d_L$

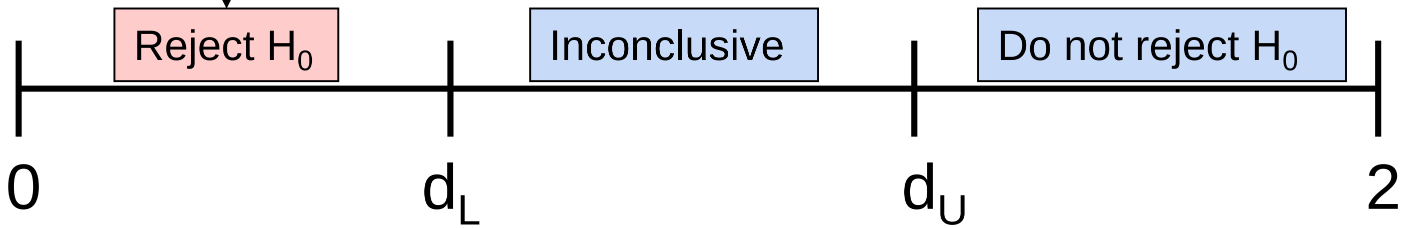

Testing for Positive Autocorrelation

(continued)

DCOVA

- Suppose we have the following time series data:

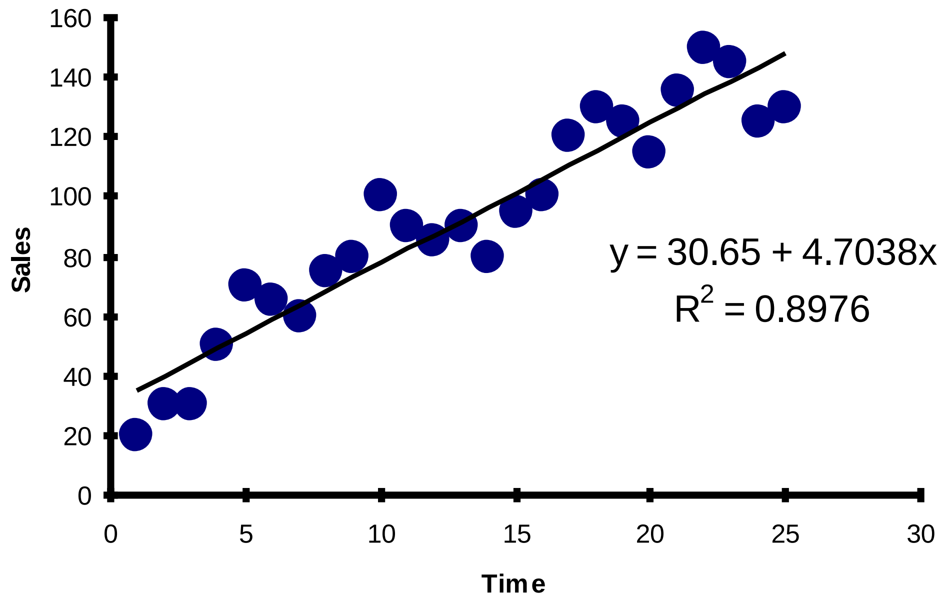

- Is there autocorrelation?

Testing for Positive Autocorrelation

(continued)

DCOVA

- Example with n = 25:

Excel/PHStat output:

Durbin-Watson Calculations	
Sum of Squared Difference of Residuals	3296.18
Sum of Squared Residuals	3279.98
Durbin-Watson Statistic	**1.00494**

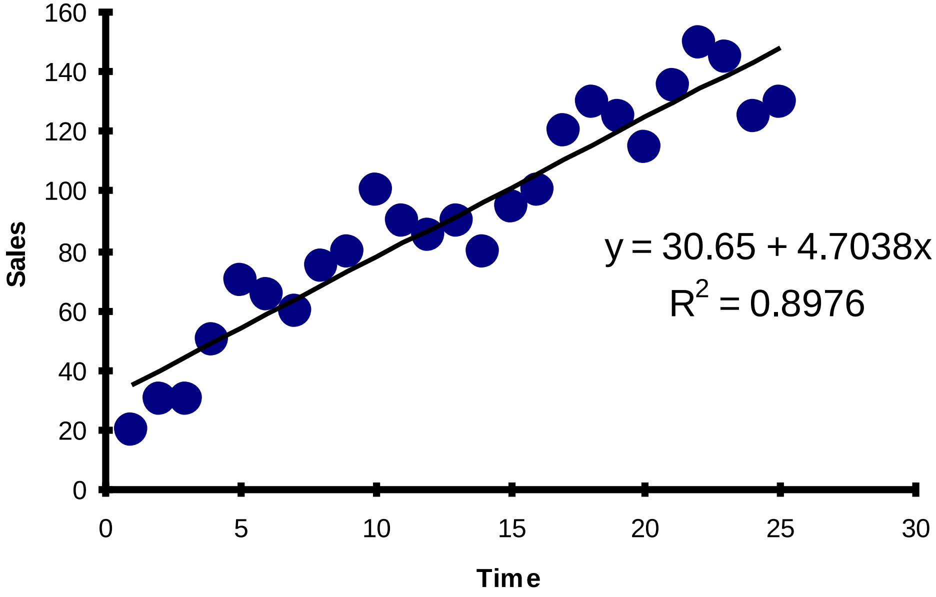

$$D = \frac{\sum_{i=2}^{n}(e_i - e_{i-1})^2}{\sum_{i=1}^{n} e_i^2} = \frac{3296.18}{3279.98} = 1.00494$$

Testing for Positive Autocorrelation

(continued)

DCOVA

- Here, n = 25 and there is k = 1 one independent variable
- Using the Durbin-Watson table, $d_L = 1.29$ and $d_U = 1.45$
- $D = 1.00494 < d_L = 1.29$, so reject H_0 and conclude that significant positive autocorrelation exists

Decision: **reject H_0** since

$D = 1.00494 < d_L$

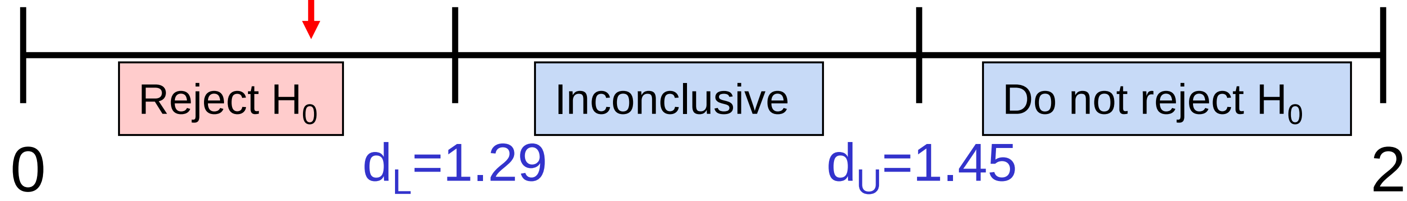

Inferences About the Slope

DCOVA

- The standard error of the regression slope coefficient (b_1) is estimated by

$$S_{b_1} = \frac{S_{YX}}{\sqrt{SSX}} = \frac{S_{YX}}{\sqrt{\sum (X_i - \bar{X})^2}}$$

where:

S_{b_1} = Estimate of the standard error of the slope

$S_{YX} = \sqrt{\frac{SSE}{n-2}}$ = Standard error of the estimate

Inferences About the Slope: t Test

DCOVA

- t test for a population slope
  - Is there a linear relationship between X and Y?
- Null and alternative hypotheses
  - H_0: $\beta_1 = 0$ (no linear relationship)
  - H_1: $\beta_1 \neq 0$ (linear relationship does exist)
- Test statistic

$$t_{STAT} = \frac{b_1 - \beta_1}{S_{b_1}}$$

$$d.f. = n - 2$$

where:

b_1 = regression slope coefficient

β_1 = hypothesized slope

S_{b1} = standard error of the slope

Inferences About the Slope: t Test Example

DCOVA

House Price in \$1000s (y)	Square Feet (x)
245	1400
312	1600
279	1700
308	1875
199	1100
219	1550
405	2350
324	2450
319	1425
255	1700

Estimated Regression Equation:

$$\text{house price} = 98.25 + 0.1098\ (\text{sq.ft.})$$

The slope of this model is 0.1098

Is there a relationship between the square footage of the house and its sales price?

Inferences About the Slope: t Test Example

DCOVA

H_0: $\beta_1 = 0$

H_1: $\beta_1 \neq 0$

From Minitab output:

Predictor	Coef	SE Coef	T	P
Constant	98.25	58.03	1.69	0.129
Square Feet	0.10977	0.03297	3.33	0.010

b_1

S_{b_1}

$$t_{STAT} = \frac{b_1 - \beta_1}{S_{b_1}} = \frac{0.10977 - 0}{0.03297} = 3.32938$$

Inferences About the Slope: t Test Example

DCOVA

Test Statistic: $t_{STAT} = 3.329$

$H_0: \beta_1 = 0$
$H_1: \beta_1 \neq 0$

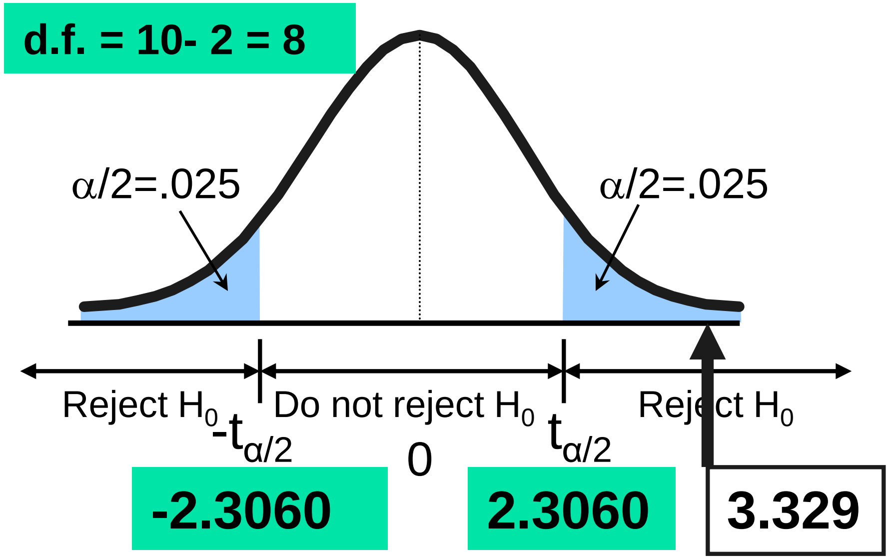

Decision: Reject H_0

There is sufficient evidence that square footage affects house price

Inferences About the Slope: t Test Example

DCOVA

H_0: $\beta_1 = 0$

H_1: $\beta_1 \neq 0$

From Minitab output:

Predictor	Coef	SE Coef	T	P
Constant	98.25	58.03	1.69	0.129
Square Feet	0.10977	0.03297	3.33	0.010

p-value

Decision: Reject H_0, since p-value < α

There is sufficient evidence that square footage affects house price.

F Test for Significance

DCOVA

- F Test statistic: $$F_{STAT} = \frac{MSR}{MSE}$$

where

$$\text{MSR} = \frac{\text{SSR}}{\text{k}}$$

$$\text{MSE} = \frac{\text{SSE}}{\text{n} - \text{k} - 1}$$

where F_{STAT} follows an F distribution with k numerator and (n – k - 1) denominator degrees of freedom

(k = the number of independent variables in the regression model)

F-Test for Significance Minitab Output

DCOVA

Analysis of Variance

Source	DF	SS	MS	F	P
Regression	1	18935	18935	11.08	0.010
Residual Error	8	13666	1708		
Total	9	32600			

p-value for the F-Test

With 1 and 8 degrees of freedom

$$F_{STAT} = \frac{MSR}{MSE} = \frac{18934.9348}{1708.1957} = 11.0848$$

F Test for Significance

(continued)

DCOVA

H_0: $\beta_1 = 0$

H_1: $\beta_1 \neq 0$

$\alpha = .05$

$df_1 = 1$ $df_2 = 8$

Critical Value:

$F_\alpha = 5.32$

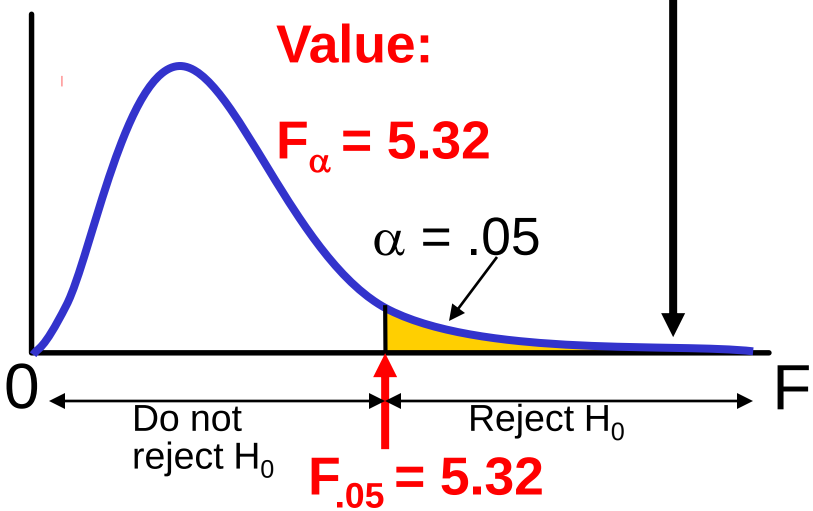

Test Statistic:

$$F_{STAT} = \frac{MSR}{MSE} = 11.08$$

Decision:

Reject H_0 at $\alpha = 0.05$

Conclusion:

There is sufficient evidence that house size affects selling price

Confidence Interval Estimate for the Slope

DCOVA

Confidence Interval Estimate of the Slope:

$$b_1 \pm t_{\alpha/2} S_{b_1}$$

d.f. = n - 2

Excel Printout for House Prices:

	Coefficients	*Standard Error*	*t Stat*	*P-value*	*Lower 95%*	*Upper 95%*
Intercept	**98.24833**	**58.03348**	**1.69296**	**0.12892**	**-35.57720**	**232.07386**
Square Feet	**0.10977**	**0.03297**	**3.32938**	**0.01039**	**0.03374**	**0.18580**

At 95% level of confidence, the confidence interval for the slope is (0.0337, 0.1858)

Confidence Interval Estimate for the Slope

(continued)

DCOVA

	Coefficients	*Standard Error*	*t Stat*	*P-value*	*Lower 95%*	*Upper 95%*
Intercept	**98.24833**	**58.03348**	**1.69296**	**0.12892**	**-35.57720**	**232.07386**
Square Feet	**0.10977**	**0.03297**	**3.32938**	**0.01039**	**0.03374**	**0.18580**

Since the units of the house price variable is \$1000s, we are 95% confident that the average impact on sales price is between \$33.74 and \$185.80 per square foot of house size

This 95% confidence interval does not include 0.

Conclusion: There is a significant relationship between house price and square feet at the .05 level of significance

Confidence Interval Estimate for the Slope from Minitab

(continued)

DCOVA

Minitab does not automatically calculate a confidence interval for the slope but provides the quantities necessary to use the confidence interval formula.

Predictor	Coef	SE Coef	T	P
Constant	98.25	58.03	1.69	0.129
Square Feet	0.10977	0.03297	3.33	0.010

$$b_1 \pm t_{\alpha/2} S_{b_1}$$

t Test for a Correlation Coefficient

DCOVA

- Hypotheses

$H_0: \rho = 0$	(no correlation between X and Y)
$H_1: \rho \neq 0$	(correlation exists)

- Test statistic

$$t_{STAT} = \frac{r - \rho}{\sqrt{\frac{1 - r^2}{n - 2}}}$$

(with n – 2 degrees of freedom)

where

$$r = +\sqrt{r^2} \quad \text{if } b_1 > 0$$

$$r = -\sqrt{r^2} \quad \text{if } b_1 < 0$$

t-test For A Correlation Coefficient

(continued)

DCOVA

Is there evidence of a linear relationship between square feet and house price at the .05 level of significance?

H_0: ρ = 0 (No correlation)

H_1: ρ ≠ 0 (correlation exists)

α =.05 , df = 10 - 2 = 8

$$t_{STAT} = \frac{r-\rho}{\sqrt{\frac{1-r^2}{n-2}}} = \frac{.762-0}{\sqrt{\frac{1-.762^2}{10-2}}} = 3.329$$

t-test For A Correlation Coefficient

(continued)

DCOVA

$$t_{STAT} = \frac{r - \rho}{\sqrt{\frac{1 - r^2}{n - 2}}} = \frac{.762 - 0}{\sqrt{\frac{1 - .762^2}{10 - 2}}} = 3.329$$

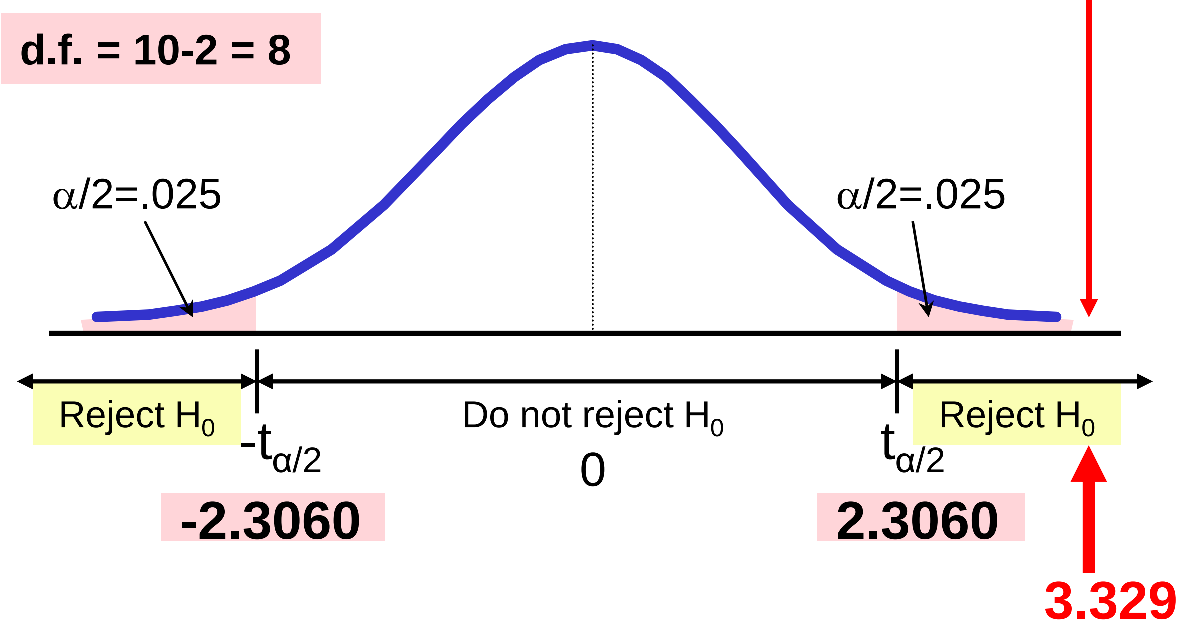

Decision:
Reject H_0

Conclusion:
There **is evidence** of a linear association at the 5% level of significance

Estimating Mean Values and Predicting Individual Values

DCOVA

Goal: Form intervals around Y to express uncertainty about the value of Y for a given X_i

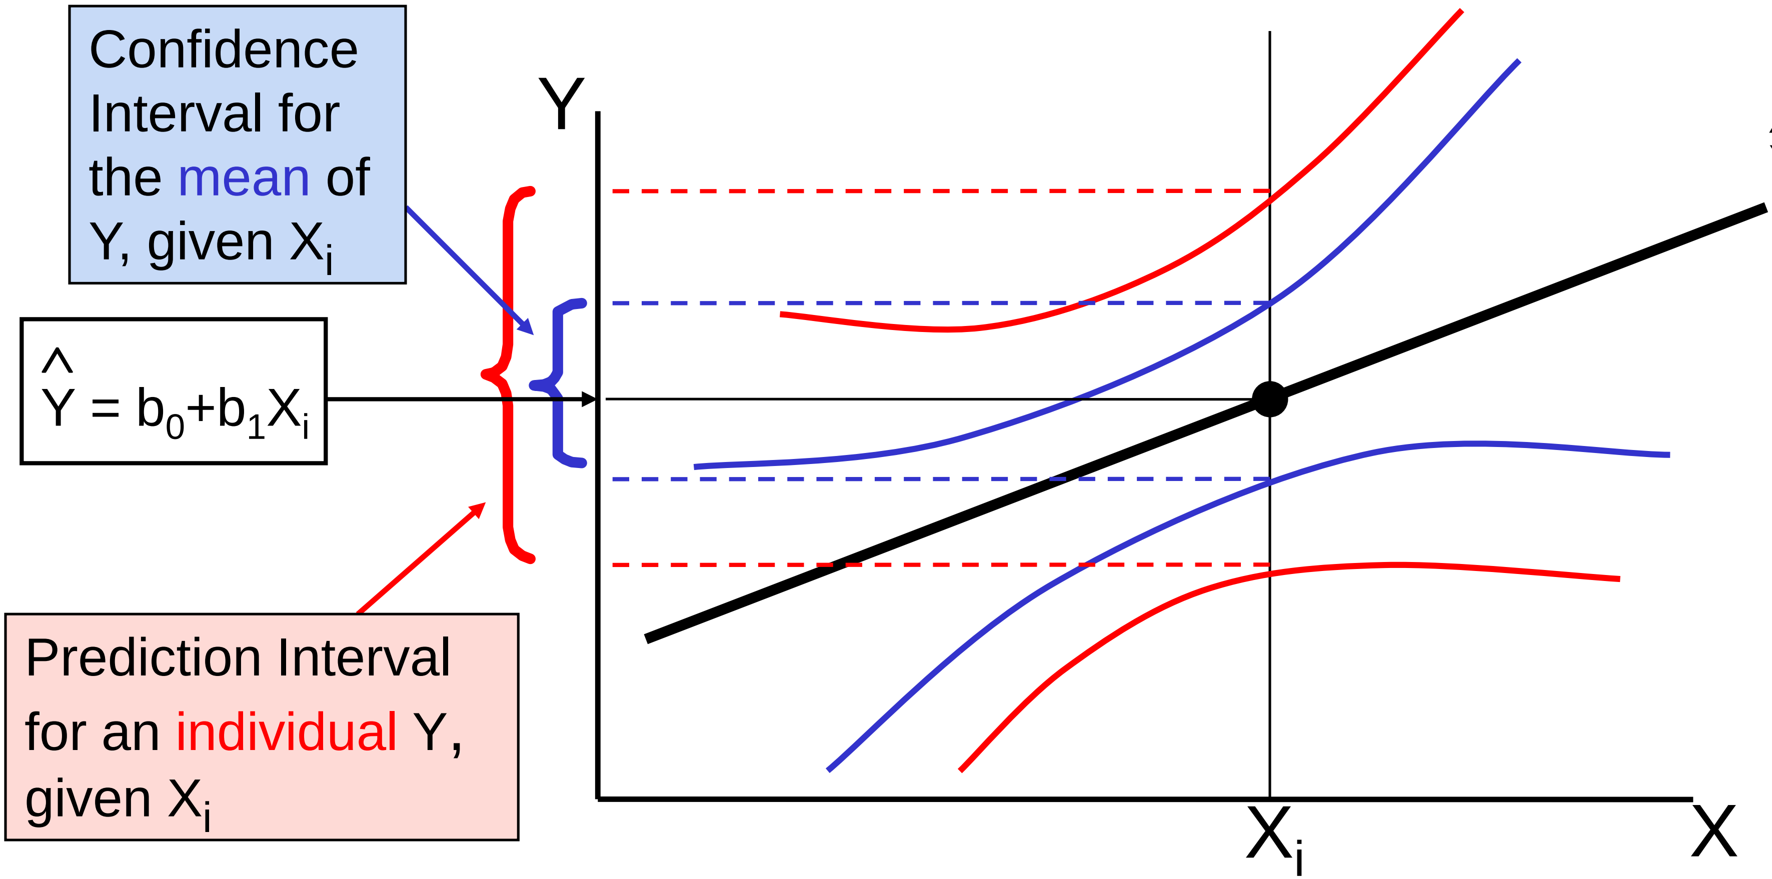

Confidence Interval for the Average Y, Given X

DCOVA

Confidence interval estimate for the **mean value of Y** given a particular X_i

Confidence interval for $\mu_{Y|X=X_i}$:

$$\hat{Y} \pm t_{\alpha/2} S_{YX} \sqrt{h_i}$$

Size of interval varies according to distance away from mean, $\bar{X}$

$$h_i = \frac{1}{n} + \frac{(X_i - \bar{X})^2}{SSX} = \frac{1}{n} + \underbrace{\frac{(X_i - \bar{X})^2}{\sum (X_i - \bar{X})^2}}$$

Prediction Interval for an Individual Y, Given X

DCOVA

Confidence interval estimate for an **Individual value of Y** given a particular X_i

$$\text{Confidence interval for } Y_{X=X_i}:$$

$$\hat{Y} \pm t_{\alpha/2} S_{YX} \sqrt{1 + h_i}$$

This extra term adds to the interval width to reflect the added uncertainty for an individual case

Estimation of Mean Values: Example

DCOVA

Confidence Interval Estimate for $\mu_{Y|X=X_i}$

Find the 95% confidence interval for the mean price of 2,000 square-foot houses

Predicted Price $\hat{Y}_i$ = 317.85 ($1,000s)

$$\hat{Y} \pm t_{0.025} S_{YX} \sqrt{\frac{1}{n} + \frac{(X_i - \bar{X})^2}{\sum (X_i - \bar{X})^2}} = 317.85 \pm 37.12$$

The confidence interval endpoints (from Excel) are 280.66 and 354.90, or from $280,660 to $354,900

Estimation of Individual Values: Example

DCOVA

Prediction Interval Estimate for $Y_{X=X_i}$

Find the 95% prediction interval for an individual house with 2,000 square feet

Predicted Price $\hat{Y}_i$ = 317.85 ($1,000s)

$$\hat{Y} \pm t_{0.025} S_{YX} \sqrt{1 + \frac{1}{n} + \frac{(X_i - \bar{X})^2}{\sum (X_i - \bar{X})^2}} = 317.85 \pm 102.28$$

The prediction interval endpoints from Excel are 215.50 and 420.07, or from $215,500 to $420,070

Finding Confidence and Prediction Intervals in Minitab

DCOVA

Confidence Interval Estimate for $\mu_{Y|X=Xi}$

Predicted Values for New Observations

```
New
Obs     Fit  SE Fit       95% CI            95% PI
  1  317.8    16.1  (280.7, 354.9)  (215.5, 420.1)
```

$\hat{Y}$

Values of Predictors for New Observations

```
New  Square
Obs    Feet
  1     2000
```

Input value(s)

Prediction Interval Estimate for $Y_{X=Xi}$

Pitfalls of Regression Analysis

- Lacking an awareness of the assumptions underlying least-squares regression
- Not knowing how to evaluate the assumptions
- Not knowing the alternatives to least-squares regression if a particular assumption is violated
- Using a regression model without knowledge of the subject matter
- Extrapolating outside the relevant range

Strategies for Avoiding the Pitfalls of Regression

- Start with a scatter plot of X vs. Y to observe possible relationship
- Perform residual analysis to check the assumptions
  - Plot the residuals vs. X to check for violations of assumptions such as homoscedasticity
  - Use a histogram, stem-and-leaf display, boxplot, or normal probability plot of the residuals to uncover possible non-normality

Strategies for Avoiding the Pitfalls of Regression

(continued)

- If there is violation of any assumption, use alternative methods or models
- If there is no evidence of assumption violation, then test for the significance of the regression coefficients and construct confidence intervals and prediction intervals
- Avoid making predictions or forecasts outside the relevant range

Chapter Summary

In this chapter we discussed

- Types of regression models
- The assumptions of regression and correlation
- Determining the simple linear regression equation
- Measures of variation
- Residual analysis
- Measuring autocorrelation

Chapter Summary

(continued)

- Making inferences about the slope
- Correlation -- measuring the strength of the association
- The estimation of mean values and prediction of individual values
- Possible pitfalls in regression and recommended strategies to avoid them